



# UNIT 5

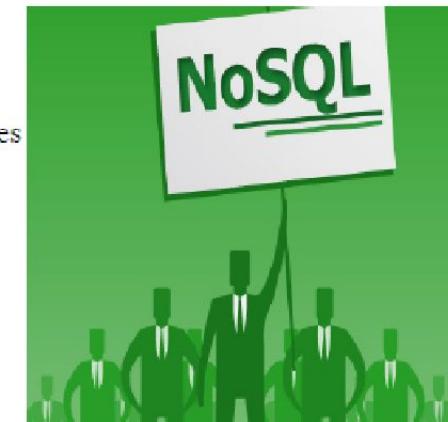
## NoSql

*Dr. G. Shobha  
Professor, CSE Department  
RV College of Engineering, Bengaluru - 59*

## NoSQL



- Term appeared in the late 90s
  - open-source relational database [Strozzi NoSQL]
  - tables as ASCII files, without SQL
- Current interpretation
  - June 11, 2009: meetup in San Francisco
  - Open-source, distributed, non-relational databases
  - Hashtag chosen: #NoSQL
  - Main features:
    - Not using SQL and the relational model
    - Open-source projects (mostly)
    - Running on clusters
    - Schemaless
  - However, no accepted precise definitions
- Most people say that NoSQL means "Not Only SQL"





## Key Points

- Relational databases have been a successful technology for twenty years, providing persistence, concurrency control, and an integration mechanism
- Application developers have been frustrated with the impedance mismatch between the relational model and the in-memory data structures
- There is a movement away from using databases as integration points towards encapsulating databases within applications and integrating through services
- The vital factor for a change in data storage was the need to support large volumes of data by running on clusters. Relational databases are not designed to run efficiently on clusters.
- NoSQL is an accidental neologism. There is no prescriptive definition—all you can make is an observation of common characteristics.
- The common characteristics of NoSQL databases are:
  - Not using the relational model
  - Running well on clusters
  - Open-source
  - Schemaless

### Popularity



## NoSQL Data Models

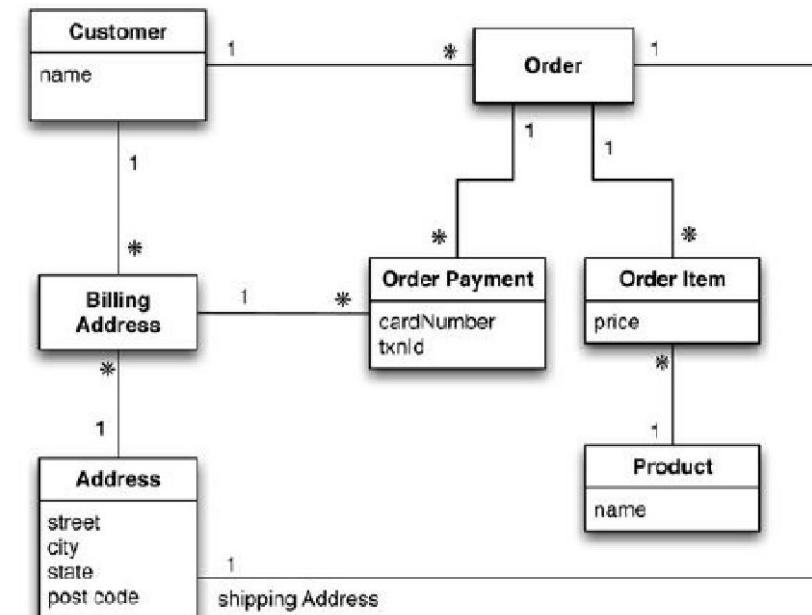
- A data model is a set of constructs for representing the information
  - Relational model: tables, columns and rows
- Storage model: how the DBMS stores and manipulates the data internally
- A data model is usually independent of the storage model
- Data models for NoSQL systems:
  - aggregate models
    - key-value,
    - document,
    - column-family
  - graph-based models



## Aggregates

- Data as atomic units that have a complex structure
  - more structure than just a set of tuples
  - example:
    - complex record with: simple fields, arrays, records nested inside
- Aggregate in Domain-Driven Design
  - a collection of related objects that we treat as a unit
  - a unit for data manipulation and management of consistency
- Advantages of aggregates:
  - easier for application programmers to work with
  - easier for database systems to handle operating on a cluster

## Example





## Relational implementation

Customer	
Id	Name
1	Martin

Orders		
Id	CustomerId	ShippingAddressId
99	1	77

Product	
Id	Name
27	NoSQL Distilled

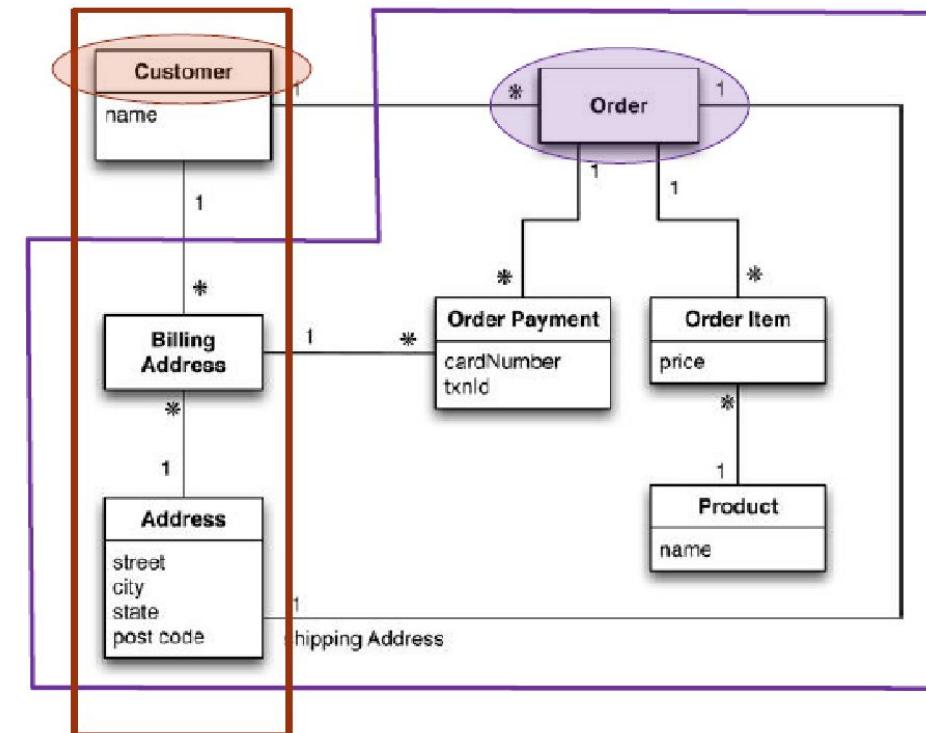
BillingAddress		
Id	CustomerId	AddressId
55	1	77

OrderItem			
Id	OrderId	ProductId	Price
100	99	27	32.45

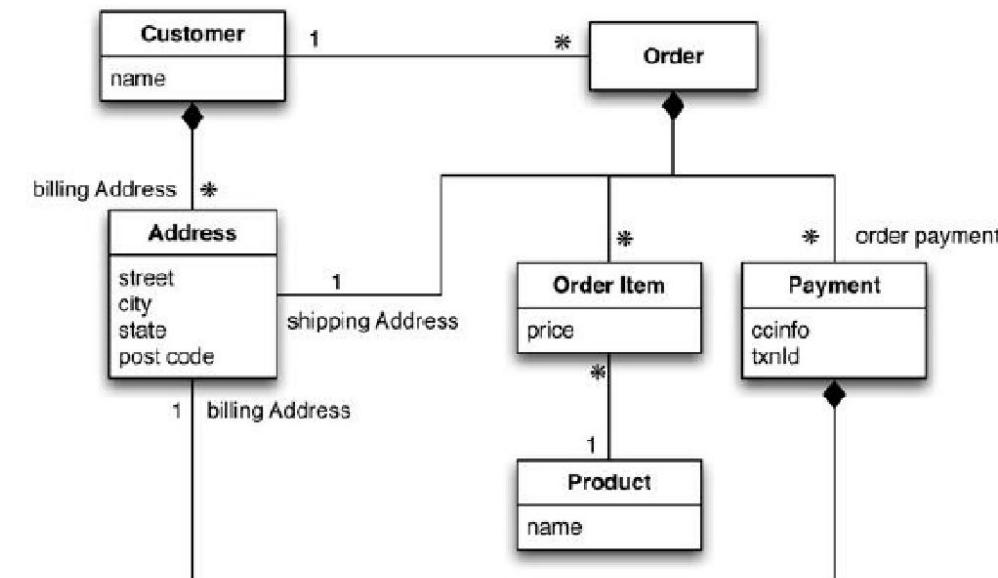
Address	
Id	City
77	Chicago

OrderPayment				
Id	OrderId	CardNumber	BillingAddressId	txmId
33	99	1000-1000	55	abelif879rft

## A possible aggregation



## Aggregate representation





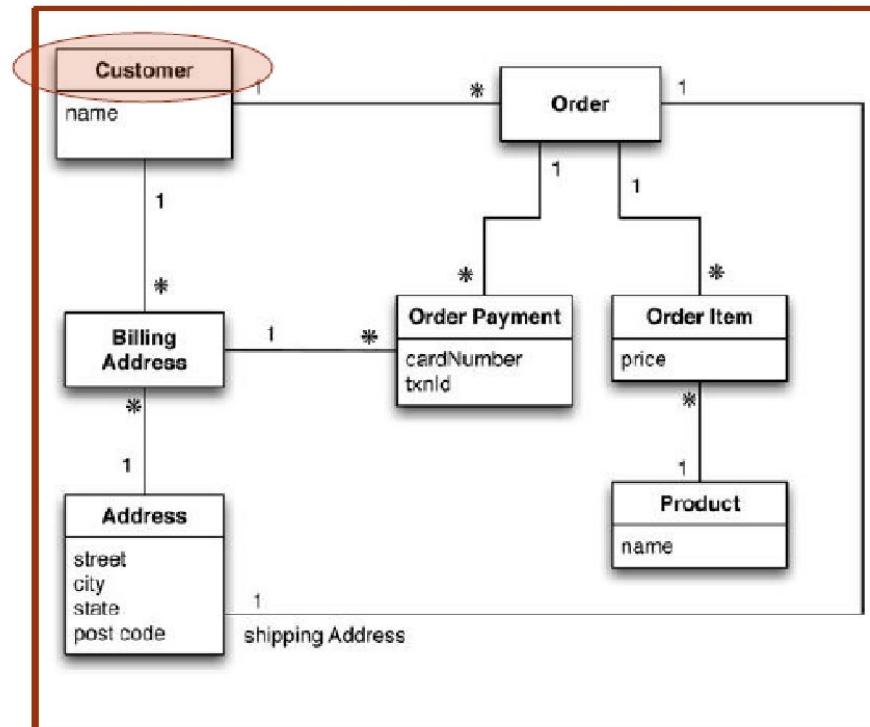
## Aggregate implementation

```
// in customers
{
  "id":1,
  "name":"Martin",
  "billingAddress":[{"city":"Chicago"}]
}

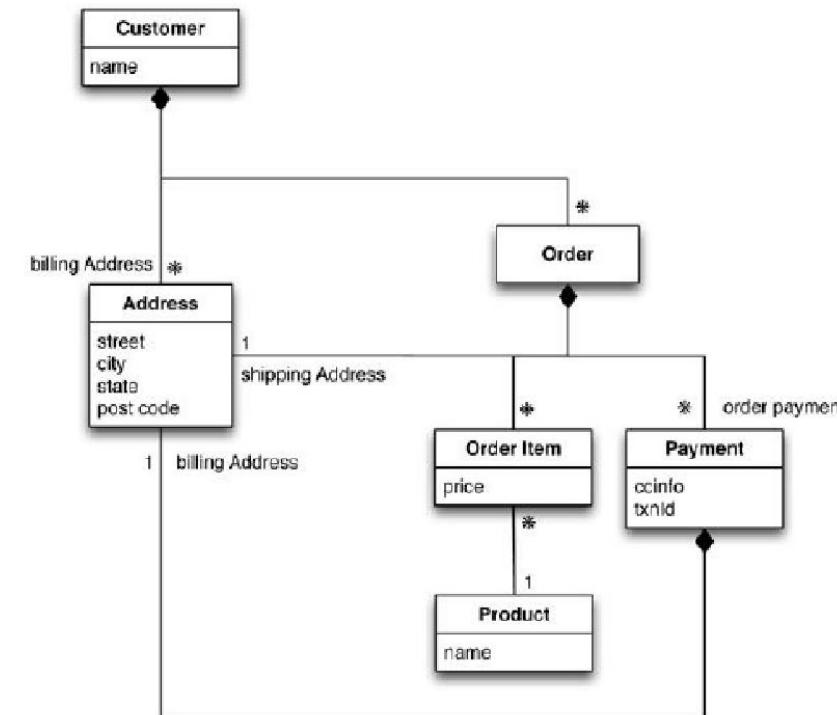
// in orders
{
  "id":99,
  "customerId":1,
  "orderItems":[
    {
      "productId":27,
      "price": 32.45,
      "productName": "NoSQL Distilled"
    }
  ],
  "shippingAddress":[{"city":"Chicago"}]
  "orderPayment":[
    {
      "ccinfo":"1000-1000-1000-1000",
      "txnid":"abelif879rft",
      "billingAddress": {"city": "Chicago"}
    }
  ],
}
```



## Another possible aggregation



## Aggregate representation (2)





## Aggregate implementation (2)

```
// in customers
{
  "customer": {
    "id": 1,
    "name": "Martin",
    "billingAddress": [{"city": "Chicago"}],
    "orders": [
      {
        "id": 99,
        "customerId": 1,
        "orderItems": [
          {
            "productId": 27,
            "price": 32.45,
            "productName": "NoSQL Distilled"
          }
        ],
        "shippingAddress": [{"city": "Chicago"}]
        "orderPayment": [
          {
            "ccinfo": "1000-1000-1000-1000",
            "txnId": "abelif879rft",
            "billingAddress": {"city": "Chicago"}
          }
        ]
      }
    ]
  }
}
```

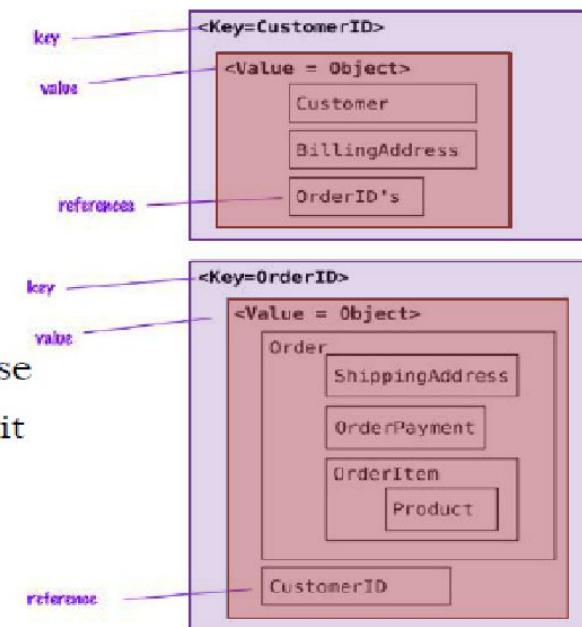
## Design strategy

- No universal answer for how to draw aggregate boundaries
- It depends entirely on how you tend to manipulate data!
  - Accesses on a single order at a time: first solution
  - Accesses on customers with all orders: second solution
- Context-specific
  - some applications will prefer one or the other
  - even within a single system
- **Focus on the unit of interaction with the data storage**
- Pros:
  - it helps greatly with running on a cluster: data will be manipulated together, and thus should live on the same node!
- Cons:
  - an aggregate structure may help with some data interactions but be an obstacle for others



## Key-Value Databases

- Strongly aggregate-oriented
  - Lots of aggregates
  - Each aggregate has a key
- Data model:
  - A set of <key,value> pairs
  - Value: an aggregate instance
- The aggregate is **opaque** to the database
  - just a big blob of mostly meaningless bits
- Access to an aggregate:
  - lookup based on its key





## Popular key-value databases



Azure Cosmos DB



## Document databases

- Strongly aggregate-oriented
  - Lots of aggregates
  - Each aggregate has a key
- Data model:
  - A set of <key,document> pairs
  - Document: an aggregate instance
- Structure of the aggregate **visible**
  - limits on what we can place in it
- Access to an aggregate:
  - queries based on the fields in the aggregate

```
# Customer object
{
  "customerId": 1,
  "name": "Martin",
  "billingAddress": [{"city": "Chicago"}],
  "payment": [
    {"type": "debit",
     "ccinfo": "1000-1000-1000-1000"}
  ]
}
```

```
# Order object
{
  "orderId": 99,
  "customerId": 1,
  "orderDate": "Nov-20-2011",
  "orderItems": [{"productId": 27, "price": 32.45}],
  "orderPayment": [{"ccinfo": "1000-1000-1000-1000",
    "txnId": "abelif879rft"}],
  "shippingAddress": {"city": "Chicago"}
}
```



## Popular document databases



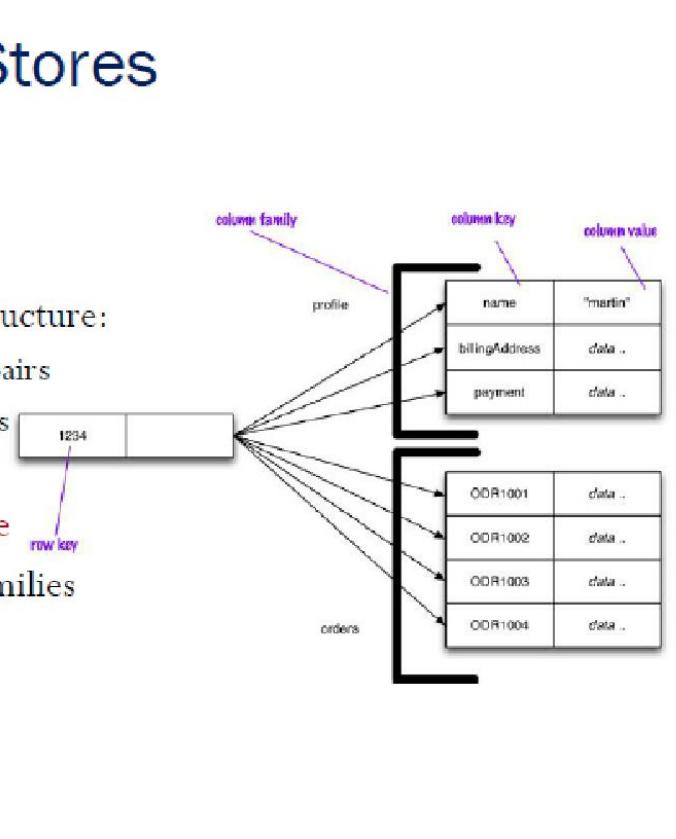
Azure Cosmos DB

## Key-Value vs Document stores

- Key-value database
  - A key plus a big blob of mostly meaningless bits
  - We can store whatever we like in the aggregate
  - We can only access an aggregate by lookup based on its key
- Document database
  - A key plus a structured aggregate
  - More flexibility in access
    - we can submit queries to the database based on the fields in the aggregate
    - we can retrieve part of the aggregate rather than the whole thing
  - Indexes based on the contents of the aggregate

## Column(-Family) Stores

- Strongly aggregate-oriented
  - Lots of aggregates
  - Each aggregate has a key
- Data model: a two-level map structure:
  - A set of <row-key, aggregate> pairs
  - Each aggregate is a group of pairs <column-key,value>
- Structure of the aggregate visible
- Columns can be organized in families
  - Data usually accessed together
- Access to an aggregate:
  - accessing the row as a whole
  - picking out a particular column





## Properties of Column Stores

- Operations also allow picking out a particular column
  - `get('1234', 'name')`
- Each column:
  - has to be part of a single column family
  - acts as unit for access
- You can add any column to any row, and rows can have very different columns
- You can model a list of items by making each item a separate column.
- Two ways to look at data:
  - Row-oriented
    - Each row is an aggregate
    - Column families represent useful chunks of data within that aggregate.
  - Column-oriented:
    - Each column family defines a record type
    - Row as the join of records in all column families

## Cassandra



- Skinny row
  - few columns
  - same columns used by many different rows
  - each row is a record and each column is a field
- Wide row
  - many columns (perhaps thousands)
  - rows having very different columns
  - models a list, with each column being one element in that list
- A column store can contain both field-like columns and list-like columns



## Popular column stores



HYPERTABLE





## Key Points

- An aggregate is a collection of data that we interact with as a unit.
- Aggregates form the boundaries for ACID operations with the database
- Key-value, document, and column-family databases can all be seen as forms of aggregate-oriented database
- Aggregates make it easier for the database to manage data storage over clusters
- Aggregate-oriented databases work best when most data interaction is done with the same aggregate
- Aggregate-ignorant databases are better when interactions use data organized in many different formations

## Relationships

- Relationship between different aggregates:
  - Put the ID of one aggregate within the data of the other
  - Join: write a program that uses the ID to link data
  - The database is ignorant of the relationship in the data

```
// in orders
{
  "id":99,
  "customerId":1,
  "orderItems":[
    {
      "productId":27,
      "price": 32.45,
      "productName": "NoSQL Distilled"
    }
  ],
  "shippingAddress":[{"city":"Chicago"}]
  "orderPayment":[
    {
      "ccinfo":"1000-1000-1000-1000",
      "txnid":"abelif879rft",
      "billingAddress": {"city": "Chicago"}
    }
  ],
}

// in customers
{
  "id":1,
  "name":"Martin",
  "billingAddress":[{"city":"Chicago"}]
}
```



## Relationship management

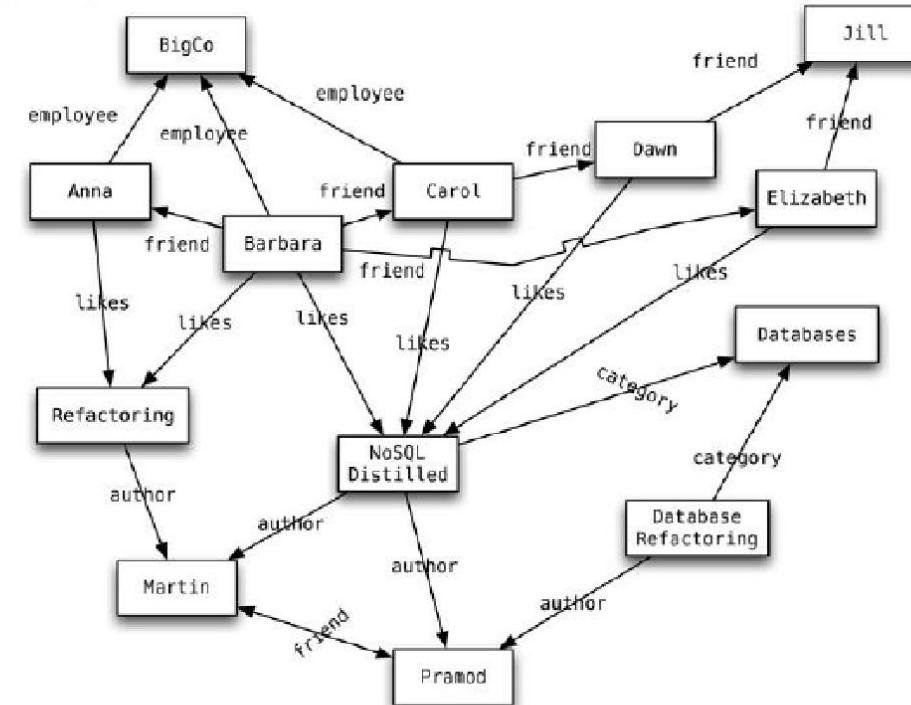
- Many NoSQL databases provide ways to make relationships visible to the database
  - Document stores makes use of indexes
  - Riak (key-value store) allows you to put link information in metadata
- But what about updates?
  - Aggregate-oriented databases treat the aggregate as the unit of data-retrieval.
  - Atomicity is only supported within the contents of a single aggregate.
  - Updates over multiple aggregates at once is a programmer's responsibility!
  - In contrast, relational databases provide ACID guarantees while altering many rows through transactions



## Graph Databases

- Graph databases are motivated by a different frustration with relational databases
  - Complex relationships require complex join
- Goal:
  - Capture data consisting of complex relationships
  - Data naturally modelled as graphs
  - Examples: Social networks, Web data, maps, networks.

## A graph database



Possible query: “find the authors of books in the Databases category that a friend of mine likes.”



## Popular graph databases



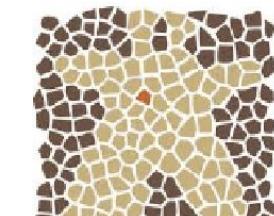
TITAN



AllegroGraph  
Franz Inc.



8





## Data model of graph databases

- Basic characteristic: nodes connected by edges (also called arcs).
- Beyond this: a lot of variation in data models
  - Neo4J stores Java objects to nodes and edges in a schemaless fashion
  - InfiniteGraph stores Java objects, which are subclasses of built-in types, as nodes and edges.
  - FlockDB is simply nodes and edges with no mechanism for additional attributes
- Queries
  - Navigation through the network of edges
  - You do need a starting place
  - Nodes can be indexed by an attribute such as ID.



## Graph vs Relational databases

- Relational databases
  - implement relationships using foreign keys
  - joins require to navigate around and can get quite expensive
- Graph databases
  - make traversal along the relationships very cheap
  - performance is better for highly connected data
  - shift most of the work from query time to insert time
  - good when querying performance is more important than insert speed



## Graph vs Aggregate-oriented databases

- Very different data models
- Aggregate-oriented databases
  - distributed across clusters
  - simple query languages
  - no ACID guarantees
- Graph databases
  - more likely to run on a single server
  - graph-based query languages
  - transactions maintain consistency over multiple nodes and edges



## Schemaless Databases

- Key-value store allows you to store any data you like under a key
- Document databases make no restrictions on the structure of the documents you store
- Column-family databases allow you to store any data under any column you like
- Graph databases allow you to freely add new edges and freely add properties to nodes and edges as you wish



## Pros and cons of schemaless data

- Pros:

- More freedom and flexibility
- You can easily change your data organization
- You can deal with non-uniform data

- Cons:

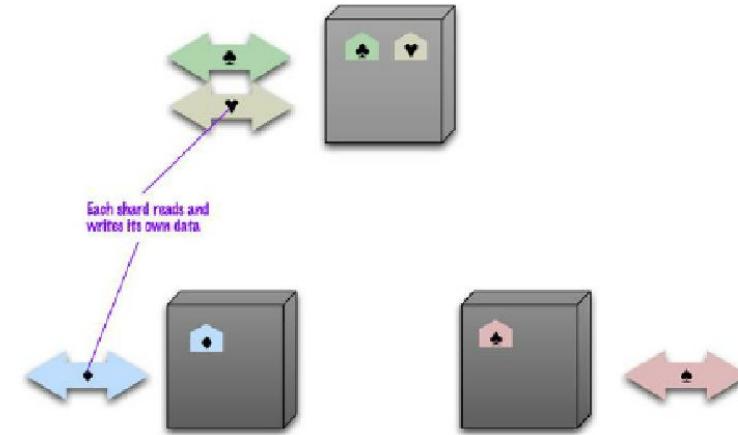
- A program that accesses data:
  - almost always relies on some form of implicit schema
  - it assumes that certain fields are present
- The implicit schema is shifted into the application code that accesses data
  - To understand what data is present you have look at the application code
- The schema cannot be used to:
  - decide how to store and retrieve data efficiently
  - ensure data consistency
- Problems if multiple applications, developed by different people, access the same database.



## Data distribution

- NoSQL systems: data distributed over large clusters
- Aggregate is a natural unit to use for data distribution
- Data distribution models:
  - Single server (is an option for some applications)
  - Multiple servers
- Orthogonal aspects of data distribution:
  - Sharding: different data on different nodes
  - Replication: the same data copied over multiple nodes

## Sharding



- Different parts of the data onto different servers
  - Horizontal scalability
  - Ideal case: different users all talking to different server nodes
  - Data accessed together on the same node — aggregate unit!
- Pros: it can improve both reads and writes
- Cons: Clusters use less reliable machines — resilience decreases



## Improving performance

Main rules of sharding:

1. Place the data close to where it's accessed
  - Orders for Boston: data in your eastern US data center
2. Try to keep the load even
  - All nodes should get equal amounts of the load
3. Put together data that may be read in sequence
  - Same order, same node
  - Many NoSQL databases offer **auto-sharding**
    - the database takes on the responsibility of sharding



## Replication



It comes in two forms:



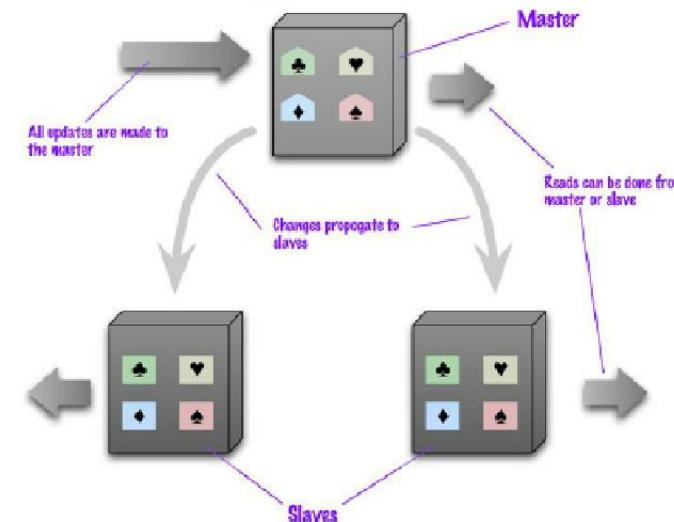
master-slave



peer-to-peer



## Master-Slave Replication



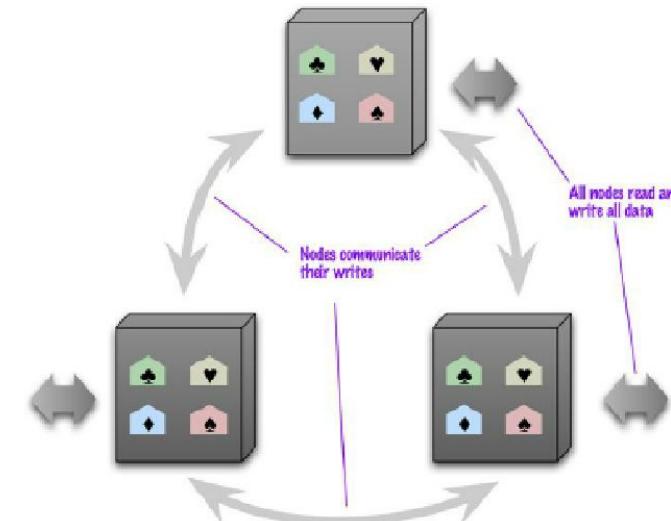
- Master
  - is the authoritative source for the data
  - is responsible for processing any updates to that data
  - can be appointed manually or automatically
- Slaves
  - A replication process synchronizes the slaves with the master
  - After a failure of the master, a slave can be appointed as new master very quickly



## Pros and cons of Master-Slave Replication

- Pros
  - More read requests:
    - Add more slave nodes
    - Ensure that all read requests are routed to the slaves
  - Should the master fail, the slaves can still handle read requests
  - Good for datasets with a read-intensive dataset
- Cons
  - The master is a bottleneck
    - Limited by its ability to process updates and to pass those updates on
    - Its failure does eliminate the ability to handle writes until:
      - the master is restored or
      - a new master is appointed
  - Inconsistency due to slow propagation of changes to the slaves
  - Bad for datasets with heavy write traffic

## Peer-to-Peer Replication



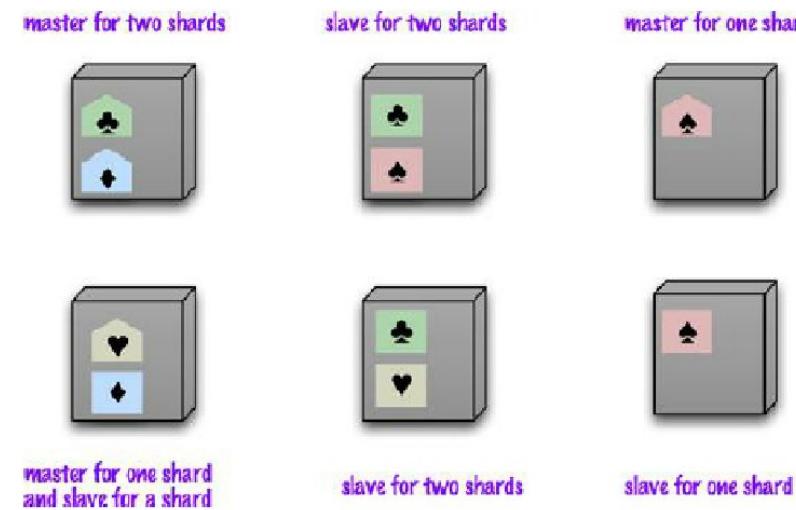
- All the replicas have equal weight, they can all accept writes
- The loss of any of them doesn't prevent access to the data store.



## Pros and cons of peer-to-peer replication

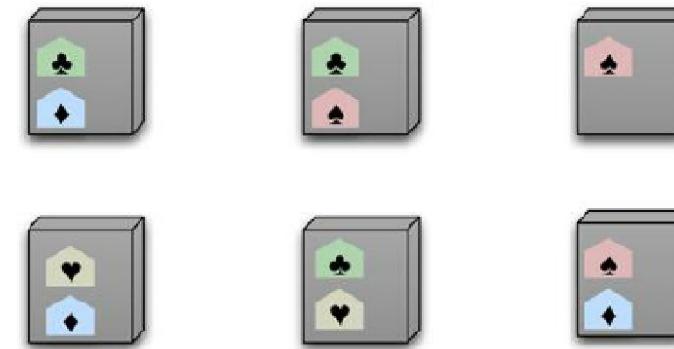
- Pros:
  - you can ride over node failures without losing access to data
  - you can easily add nodes to improve your performance
- Cons:
  - Inconsistency!
    - Slow propagation of changes to copies on different nodes
      - Inconsistencies on read lead to problems but are relatively transient
    - Two people can update different copies of the same record stored on different nodes at the same time - a **write-write conflict**.
      - Inconsistent writes are forever.

## Sharding and Replication on MS



- We have multiple masters, but each data only has a single master.
- Two schemes:
  - A node can be a master for some data and slaves for others
  - Nodes are dedicated for master or slave duties

## Sharding and Replication on P2P



- Usually each shard is present on three nodes
- A common strategy for column-family databases



## Key points

- There are two styles of distributing data:
  - Sharding distributes different data across multiple servers
    - each server acts as the single source for a subset of data.
  - Replication copies data across multiple servers
    - each bit of data can be found in multiple places.
- A system may use either or both techniques.
- Replication comes in two forms:
  - Master-slave replication makes one node the authoritative copy that handles writes while slaves synchronize with the master and may handle reads.
  - Peer-to-peer replication allows writes to any node; the nodes coordinate to synchronize their copies of the data.
- Master-slave replication reduces the chance of update conflicts but peer-to-peer replication avoids loading all writes onto a single point of failure.