

Introduction to NoSQL

- Avinash Kumar

RDBMS

- ❑ RDBMS: Relational Database Management System
- ❑ Relational Databases persist data in tabular form structured to recognize relations between stored items of information.
- ❑ DB2, MySQL, PostgreSQL, SQLite etc

Pros:

- ❑ ACID
- ❑ SQL - Structured Query Language
- ❑ Strong consistency, concurrency and recovery.

e_id	e_name	e_salary	e_age	e_gender	e_dept
1	Sam	95000	45	Male	Operations
2	Bob	80000	21	Male	Support
3	Anne	125000	25	Female	Analytics
4	Julia	73000	30	Female	Analytics
5	Matt	159000	33	Male	Sales
6	Jeff	112000	27	Male	Operations

Cons:

- ❑ Not built to be distributed.
- ❑ Joins are expensive
- ❑ Hard to scale horizontally
- ❑ Impedance mismatch

Case Analysis

Service with RDMS
for data storage



More read
operations



ACID?

Cache: Frequently
executed query



More write
operations



Vertical
scaling




New Features, Caused DB
complexity, lots of SQL joins



Denormalize
data

Expand to other region
Need distributed approach



NoSQL

- ❑ Cluster friendly
- ❑ High availability and scalability
- ❑ Schema less, non-tabular and non-relational
- ❑ Big data capabilities
- ❑ Avoids upfront schema design
- ❑ Transactions are handled at application layer
- ❑ Provides easy and frequent changes to db
- ❑ Horizontal scaling
- ❑ Fixes impedance mismatch
- ❑ Allows faster development




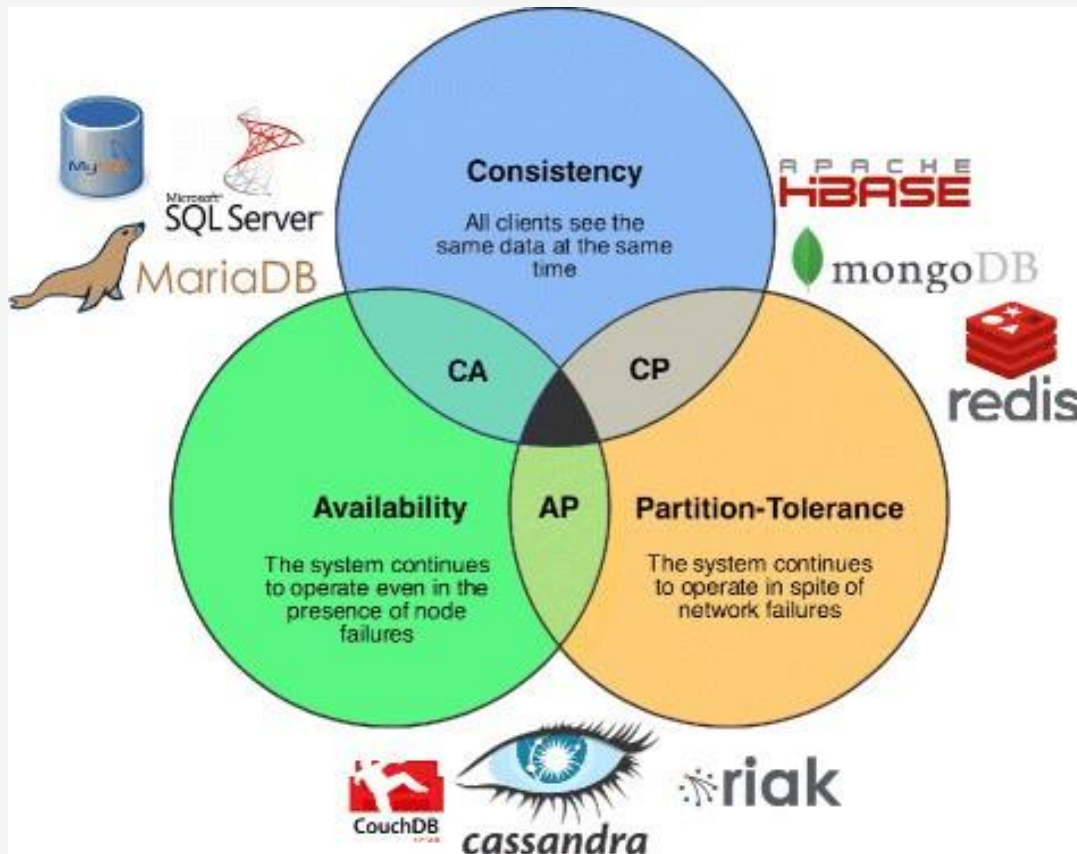
Cap Theorem

According to CAP theorem a distributed system can only provide at most 2 of these three properties

❑ Consistency

❑ Availability

❑ Partition Tolerance



BASE

- ❑ Basically available: Nodes in a distributed environment can go down but the whole system shouldn't be affected.
- ❑ Soft state (scalable): The state of the system and data could change over time.
- ❑ Eventually consistency: Given enough time, data will be consistent across the distributed system.

Aggregate Data Models

NoSQL databases are divided in four major data models:

- ☐ Key-Value
- ☐ Document
- ☐ Column family
- ☐ Graph



Key-Value

- ❑ Key serves as unique identifier
- ❑ Value can be anything ranging from simple objects to complex compound objects
- ❑ Redis, MemcacheDB, Riak etc.

Use cases:

- ❑ Storing user session data
- ❑ Maintaining schema-less profile
- ❑ Storing user preferences
- ❑ Storing shopping kart data

Key	Value
101	Simple string value
102	{ "customerId": 21 "orderItems": [{ "productId": 201 "quantity": 2 "cost": 420 }, { "productId": 201 "quantity": 2 "cost": 420 }] }

Document

- ❑ Similar to Key-value
- ❑ Difference is that, the value contains structure or semi-structured data
- ❑ Value is called document and can be in JSON, XML, BSON formats
- ❑ CouchDB, MongoDB etc.

Use cases:

- ❑ Ecommerce platforms
- ❑ Content management systems
- ❑ Analytics platform
- ❑ Blogging platform

airport_1254 Q

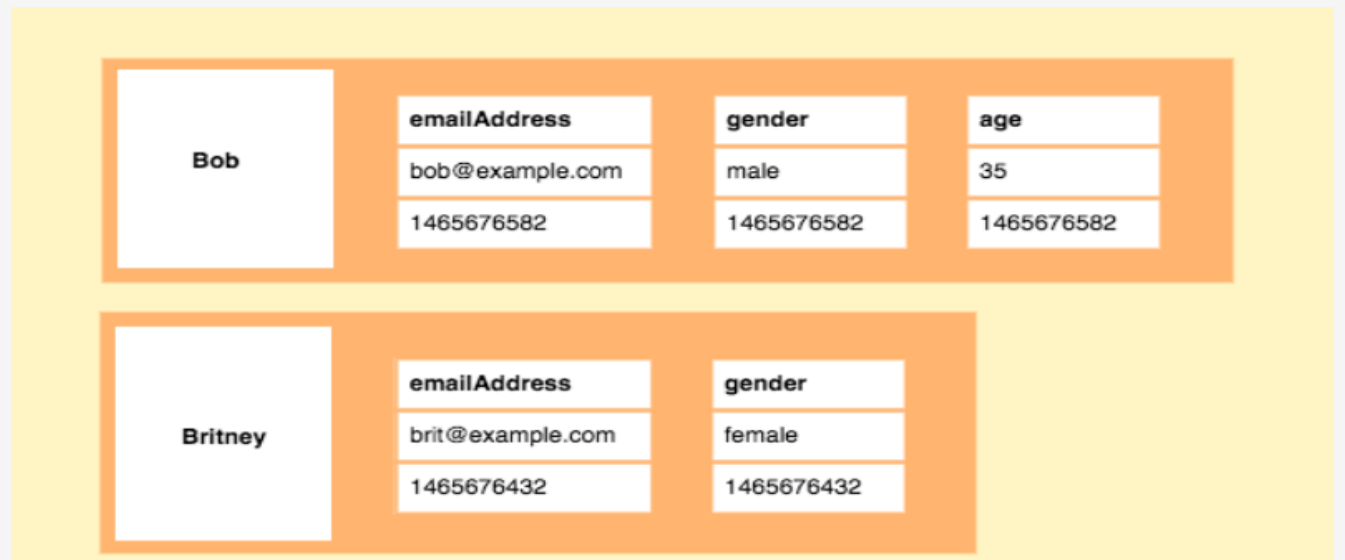
```
1 {  
2   "airportname": "Calais Dunkerque",  
3   "city": "Calais",  
4   "country": "France",  
5   "faa": "CQF",  
6   "geo": {  
7     "alt": 12,  
8     "lat": 50.962097,  
9     "lon": 1.954764  
10  },  
11   "icao": "LFAC",  
12   "id": 1254,  
13   "type": "airport",  
14   "tz": "Europe/Paris"  
15 }
```

Column family

- ❑ Data is stored in cells grouped in columns of data
- ❑ Columns are logically grouped into column families
- ❑ Column families are groups of similar data that is usually accessed together
- ❑ Cassandra, Hbase etc.

Use cases:

- ❑ Content management systems
- ❑ Systems that maintains counters
- ❑ Services having expiry usage
- ❑ Systems requiring heavy write requests
- ❑ Log aggregators

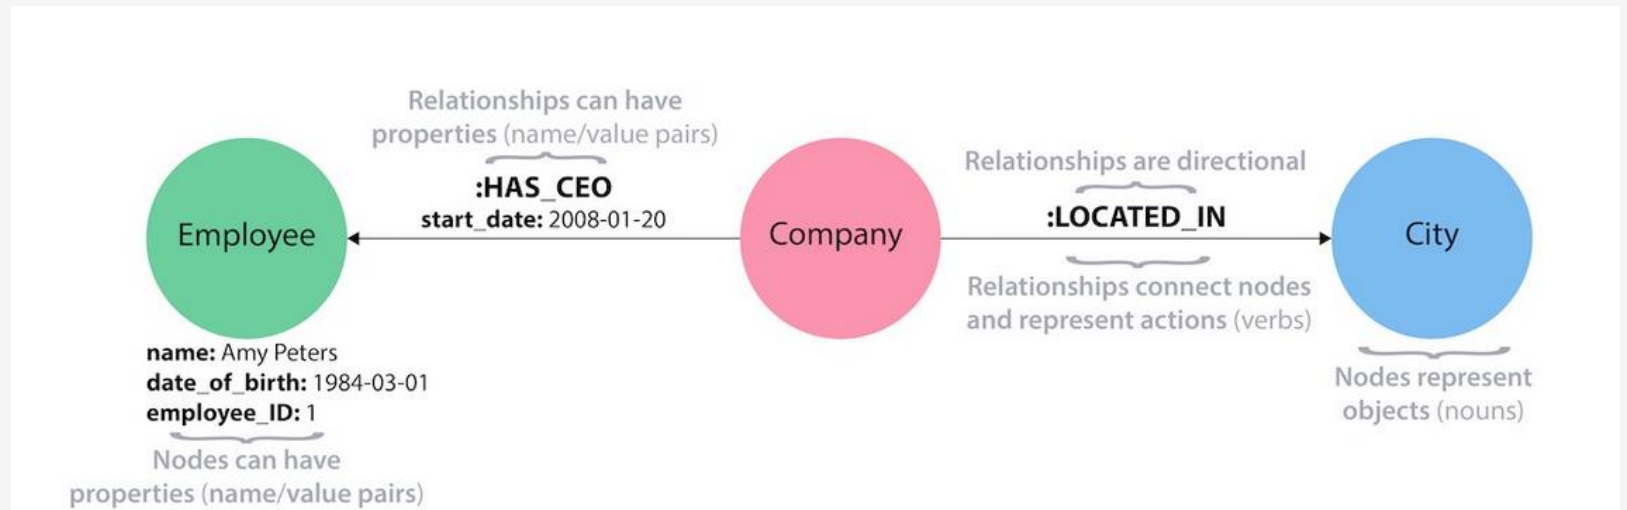


Graph

- ❑ Store entities and relationship as node and edge of the graph
- ❑ Entity contains the data and Relationship is how entity are interlinked
- ❑ Relationship can be defined on-the-fly
- ❑ Neo4j, OrientDB etc.

Use cases:

- ❑ Social network application
- ❑ Location based services
- ❑ Network and IT operations
- ❑ Fraud detection



Replication

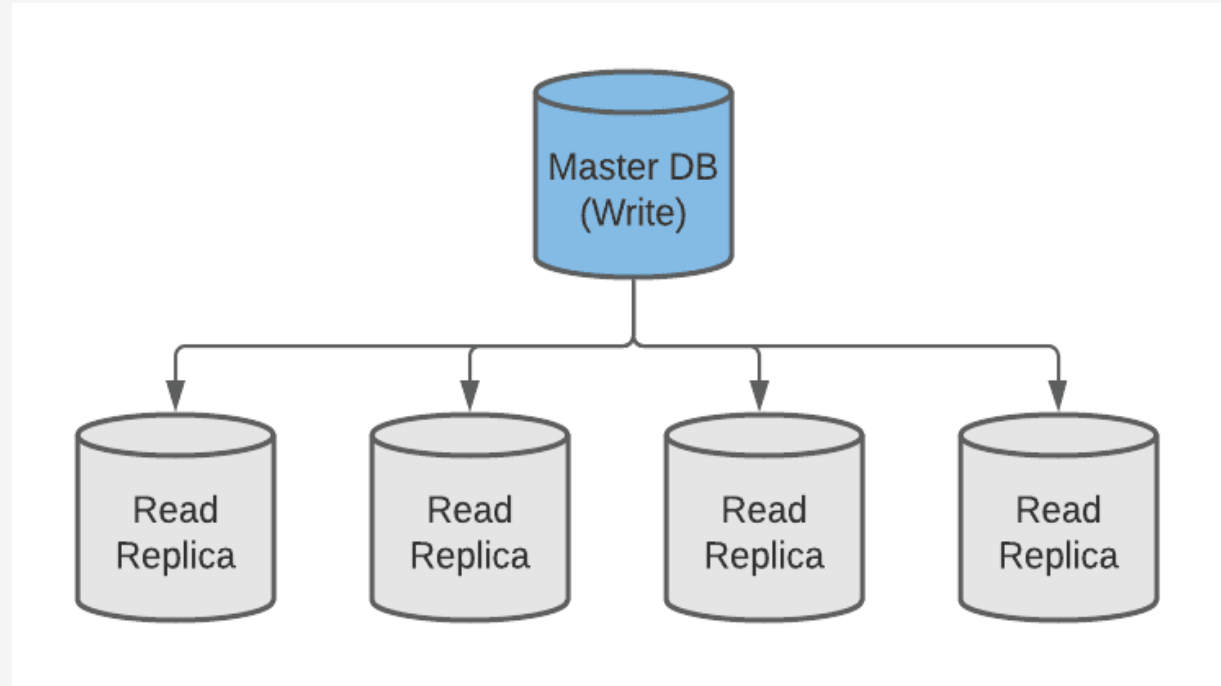
Replication copies data across multiple servers, so each bit of data can be found in multiple places.

Synchronous:

- ☐ Perform write
- ☐ Data got replicated to all nodes
- ☐ Gets back success acknowledgement

Asynchronous:


- ☐ Perform write
- ☐ Gets back success acknowledgement
- ☐ Data got replicated to all nodes in background



Type of replication


Master-Slave:

- ❑ One node is in-charge
- ❑ Write happens on master and propagates to slave nodes where read happens
- ❑ Good for read-intensive services
- ❑ Single point of failure



Peer-to-Peer:

- ❑ All nodes are equal
- ❑ Write-read can happen on any node
- ❑ Can easily withstand node failure
- ❑ Write conflicts can be the issue



Sharding


Sharding distributes different data across multiple servers, so each server acts as the single source for a subset of data.

Advantages:

- ❑ Increased Read/Write Throughput
- ❑ Increased storage capacity
- ❑ High availability

Disadvantages:

- ❑ Query overhead
- ❑ Poor shard key causes hotspots
- ❑ Increased Infrastructure Costs



Sharding Techniques

Algorithmic sharding:

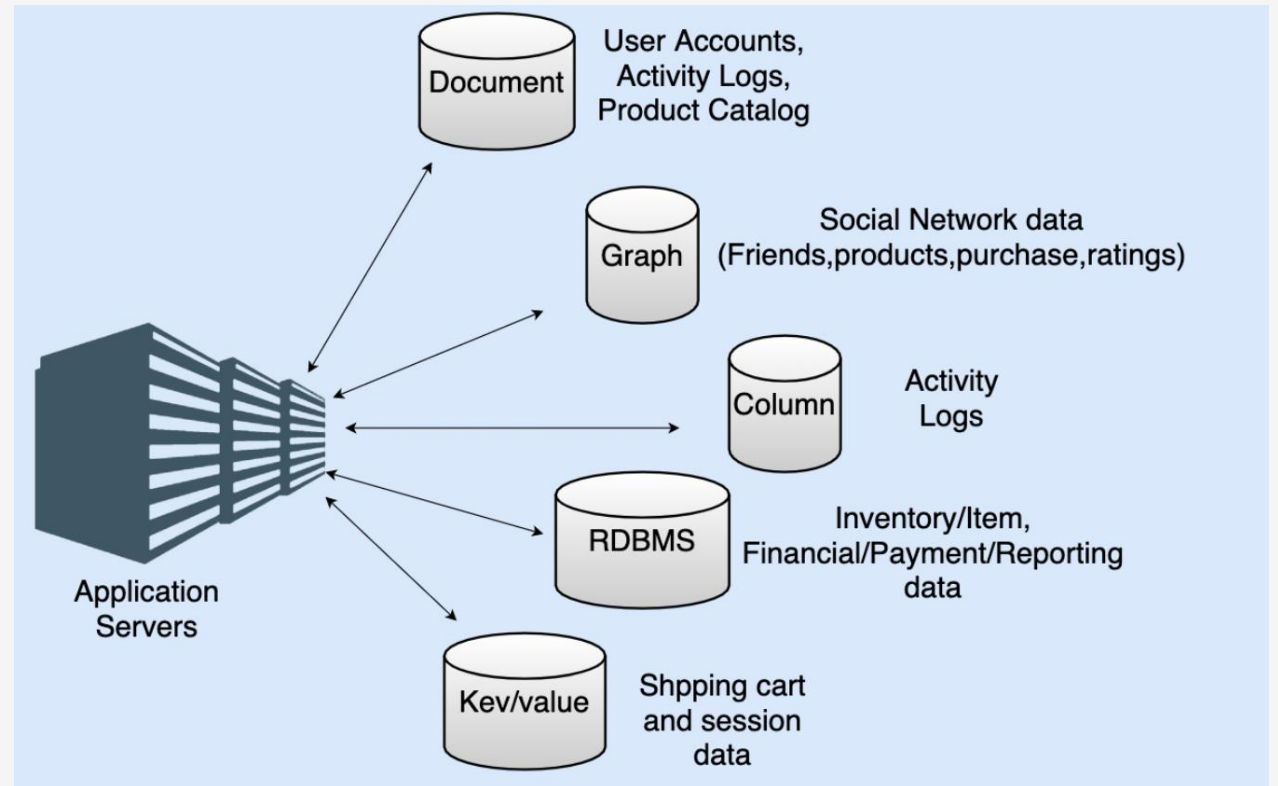
- ❑ Algorithm takes data and decides the shard it belongs
- ❑ e.g. Hash Value = $ID \% \text{Number of Shards}$
- ❑ Increasing no of shards will require rebalance

Dynamic sharding:

- ❑ Lookup service is maintained for the data and shard mapping
- ❑ e.g. Geography based sharding
- ❑ Lookup tables should be chosen wisely so that it doesn't grow very large
- ❑ Lookup service can be single point failure or bottleneck

Polyglot persistence

- ❑ Modern database solution for distributed systems are a combination of **SQL** and **NoSQL** solutions
- ❑ Modern applications leverage the ACID properties of RDBMS where consistency is key
- ❑ NoSQL is used for high availability data that gets the most hits
- ❑ A single application can interact with both SQL and NoSQL databases and also different kinds of NoSQL databases at a time.



Questions?

Thank you!